


R-course:
Machine Learning using R

Decision trees

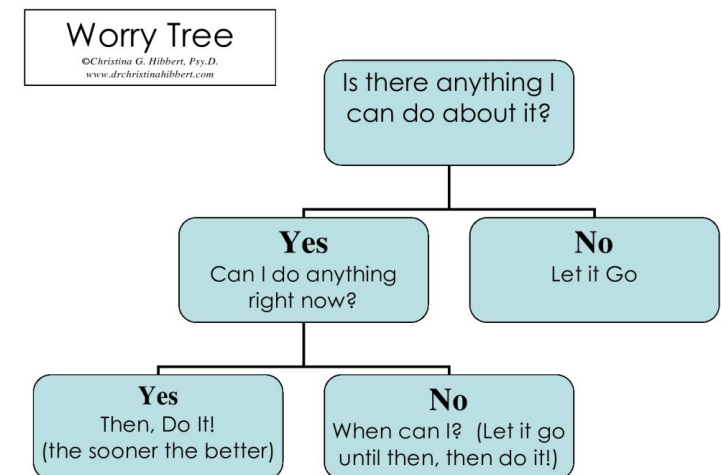


Yannick Rothacher

Zürich, 2021

What is a decision tree?

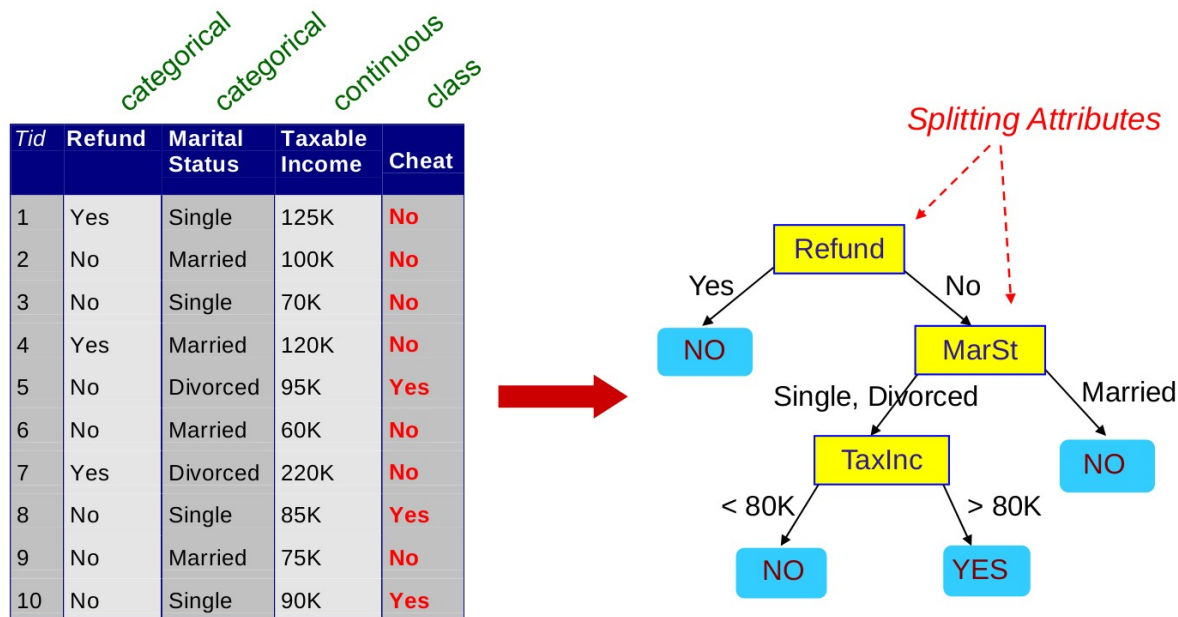
- ▶ Generally speaking, a **decision tree** is a diagram, which helps us determining a course of action (e.g. "should I worry?")



- ▶ We can use decision trees for **classification** or **regression**
 - ▶ Such trees are the result of what is referred to as **recursive partitioning**
- ▶ Let's look at an example of a decision tree used for **classification**...

Decision tree – classification

- ▶ We want to predict whether a person will cheat in his/her tax declaration based on some training data
- ▶ **"Cheat"** is the target variable (two levels: yes/no)
- ▶ The generated decision tree uses the three independent variables to model whether someone will cheat or not:



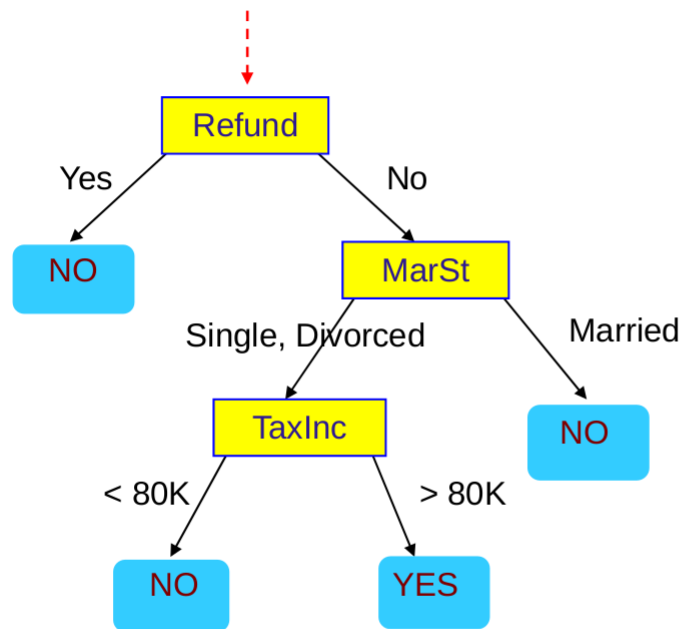
Training Data

Model: Decision Tree

Decision tree – classification

- ▶ Use the generated tree to predict outcome of a **new observation**
- ▶ Work through the tree for prediction:

Start from the root of tree.

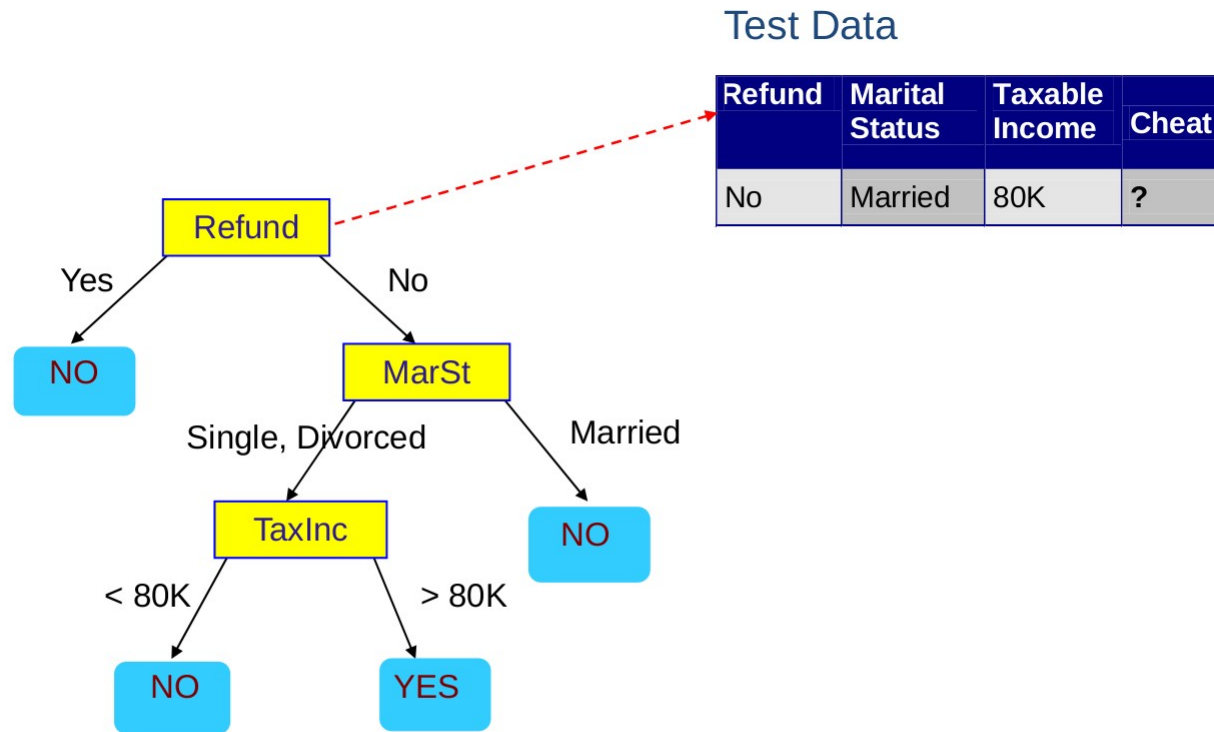


Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Decision tree – classification

- ▶ Use the generated tree to predict outcome of a **new observation**
- ▶ Work through the tree for prediction:

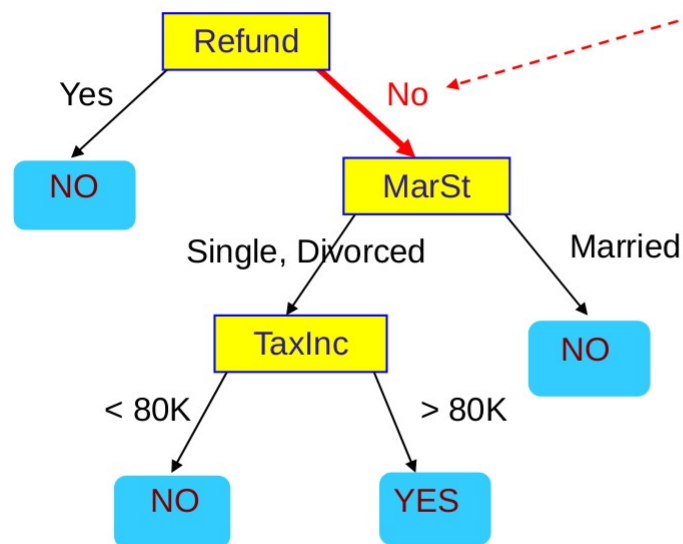


Decision tree – classification

- ▶ Use the generated tree to predict outcome of a **new observation**
- ▶ Work through the tree for prediction:

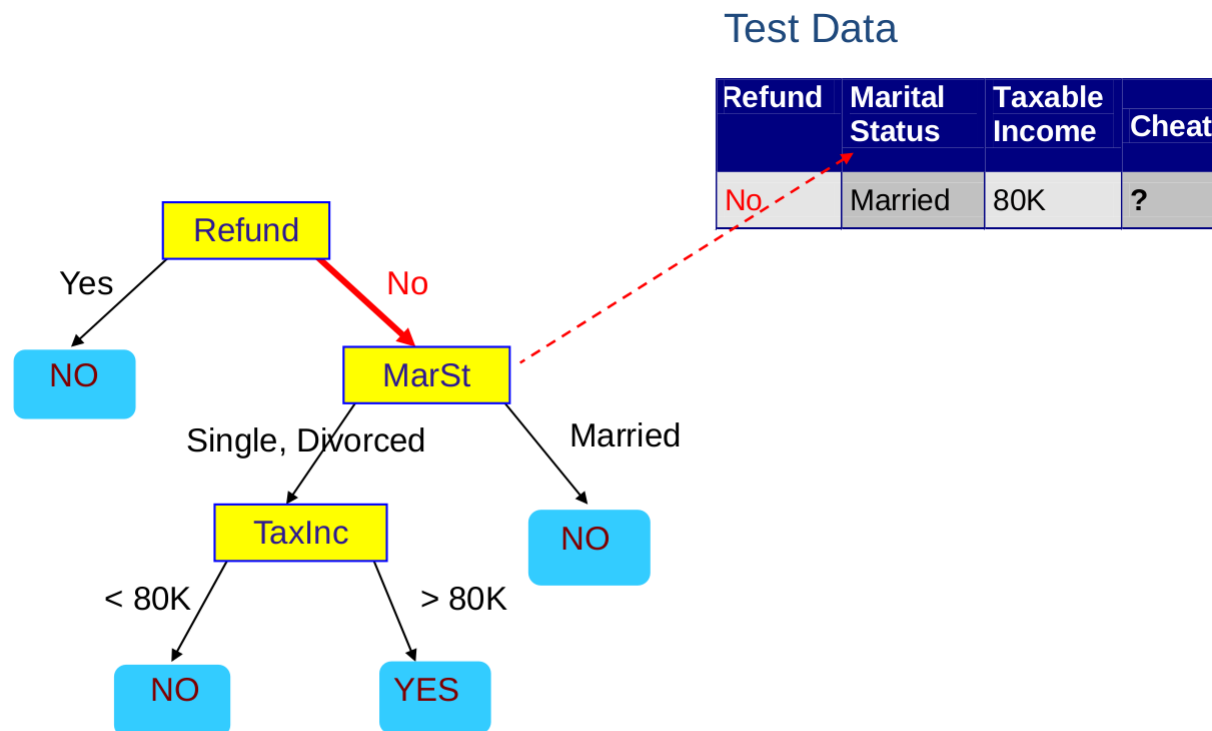
Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?




Decision tree – classification

- ▶ Use the generated tree to predict outcome of a **new observation**
- ▶ Work through the tree for prediction:



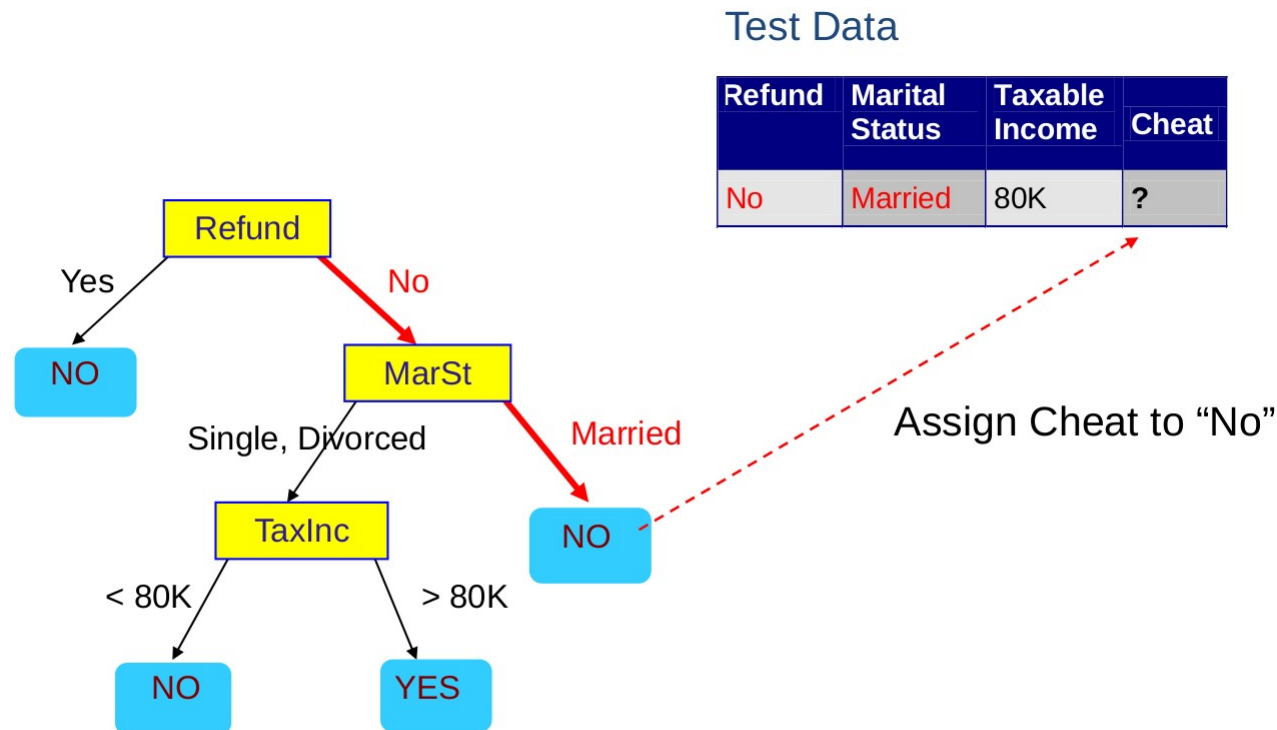
Decision tree – classification

- ▶ Use the generated tree to predict outcome of a **new observation**
- ▶ Work through the tree for prediction:



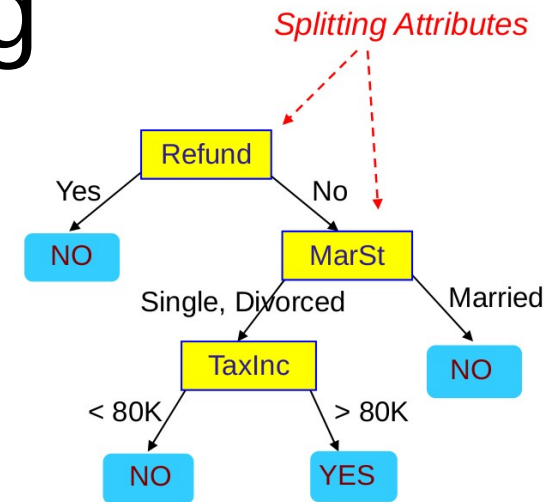
Decision tree – classification

- ▶ Use the generated tree to predict outcome of a **new observation**
- ▶ Work through the tree for prediction:






Decision tree – partitioning

- ▶ A classic decision tree splits the data at each node based on **one dependent variable**
- ▶ For the case of only **two dependent variables** we can visualize this in a 2D-coordinate system:



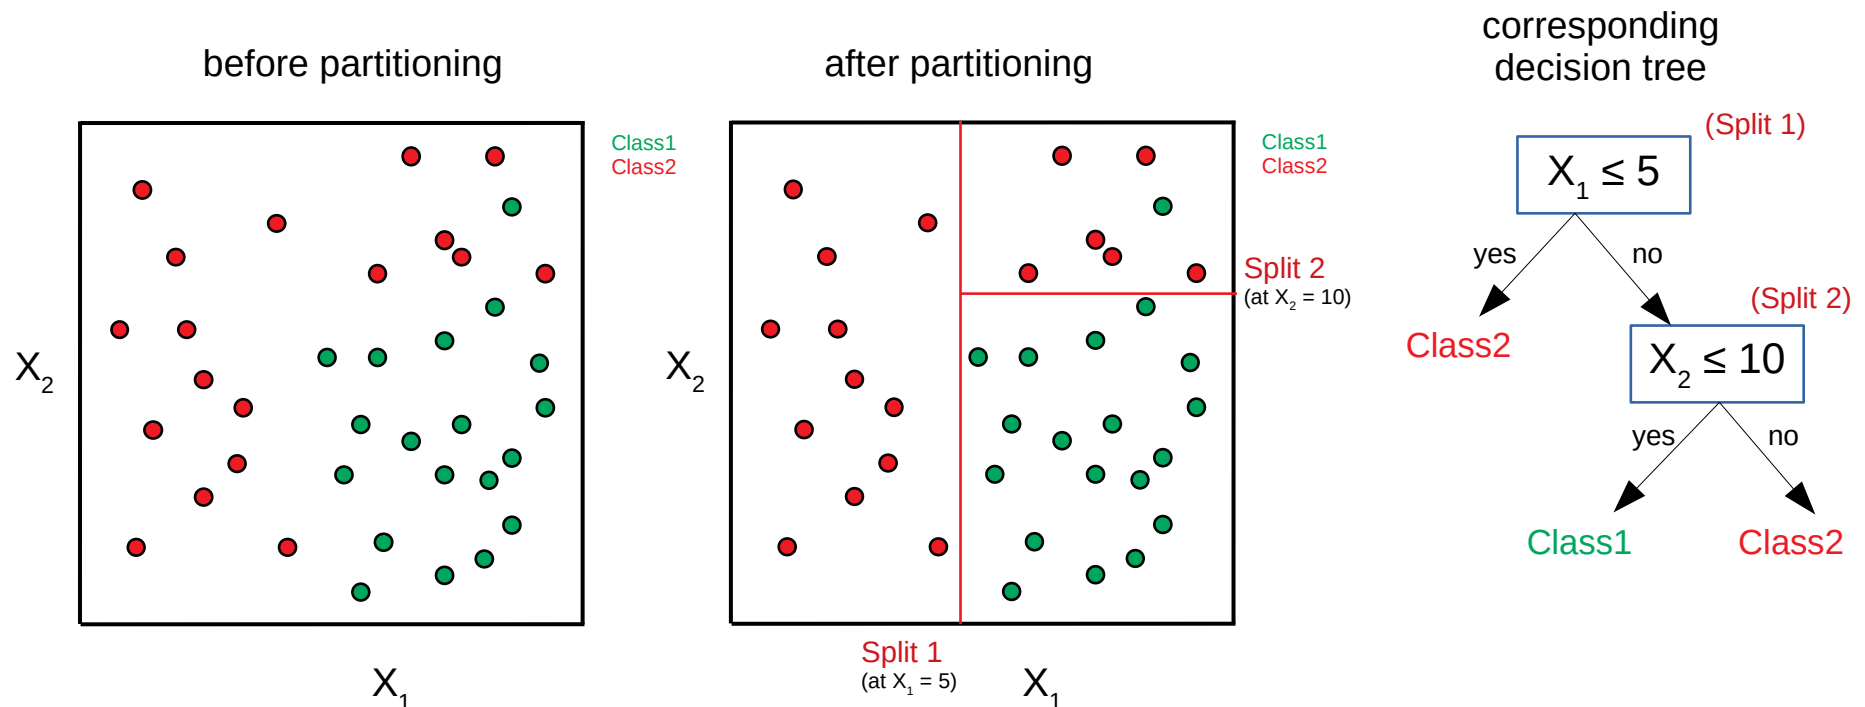
Model: Decision Tree



Only straight lines!


Decision tree – partitioning

- ▶ A classic decision tree splits the data at each node based on **one dependent variable**
- ▶ For the case of only **two dependent variables** we can visualize this in a 2D-coordinate system:





Decision tree – partitioning

- ▶ A classic decision tree splits the data at each node based on **one dependent variable**
- ▶ For the case of only **two dependent variables** we can visualize this in a 2D-coordinate system:




Decision trees – over- and underfitting

- ▶ How **deep** should a tree be allowed to grow?
 - ▶ We can always grow a tree until there are only pure partitions left (right figure below)
 - ▶ Same problem of **over- and underfitting** like with the KNN-classifier!
- ▶ Which one of the trees below is better?





Recap: Test-error

- ▶ Like in the case of the KNN-classifier we want to minimize the **test-error**
- ▶ Evaluation methods like e.g. **cross-validation** can of course also be applied to decision trees
- ▶ Prevent **over-fitting**: Fully grown trees can be "**pruned**" (cut shorter) depending on whether a branch improves the fit to data evaluated by cross-validation



How are the partitions generated?

- ▶ Building a decision tree from scratch involves the following two questions at each node:
 - ▶ Which variable should be used for the next split?
 - ▶ Where along the chosen variable should we split?
- ▶ How can we decide what the **best variable** and **split-location** is?
- ▶ We will look at two different methods to solve these questions
 - ▶ 1) classic partitioning based on **impurity measure** (e.g. Gini-index)
 - ▶ 2) bias-free partitioning based on **significance tests**



Classification trees – impurity measure


- ▶ **Main idea:** Try to find the split, which reduces the **impurity** of the data the most
- ▶ How can we measure "impurity"?
- ▶ The **Gini index** is one of the most common impurity measures

$$Gini = 1 - \sum_i^C p_i^2$$

Entropy is another common impurity measure

$$Entropy = - \sum_i^C p_i \log_2(p_i)$$

Example



$$Gini_{root} = 1 - ((11/15)^2 + (4/15)^2) = \mathbf{0.391}$$

$$Gini_L = 1 - ((2/6)^2 + (4/6)^2) = \mathbf{0.444}$$

$$Gini_R = 1 - ((0/9)^2 + (9/9)^2) = \mathbf{0}$$

Gini-reduction (weighted by number of observations in each partition):

$$Gini_decrease = Gini_{root} - (6/15) * Gini_L - (9/15) * Gini_R$$

Classification trees – impurity measure


- ▶ **Main idea:** Try to find the split, which reduces the **impurity** of the data the most
- ▶ How can we measure "impurity"?
- ▶ The **Gini index** is one of the most common impurity measures

$$Gini = 1 - \sum_i^C p_i^2$$

Entropy is another common impurity measure

$$Entropy = - \sum_i^C p_i \log_2(p_i)$$

Example



$$Gini_{root} = 1 - ((11/15)^2 + (4/15)^2) = \mathbf{0.391}$$

$$Gini_L = 1 - ((2/6)^2 + (4/6)^2) = \mathbf{0.444}$$

$$Gini_R = 1 - ((0/9)^2 + (9/9)^2) = \mathbf{0}$$

Gini-reduction (weighted by number of observations in each partition):

$$Gini_decrease = Gini_{root} - (6/15) * Gini_L - (9/15) * Gini_R$$

→ Choose the split with the largest Gini-reduction

Partition algorithm (impurity measure based)

- ▶ Start with a **single region** (encompassing all data)
- ▶ Iterate through the following steps:

For each region **R**

For each variable x_i in **R**

For each possible split s_i of x_i


Record impurity decrease

→ Choose (x_i, s_i) which gives maximum impurity decrease

→ Replace **R** with **R_R** and **R_L**

- ▶ Stop splitting either based on a **stopping rule** or when there is no more impurity reduction possible
- ▶ Stopping rule example: Stop splitting when a branch contains less than a certain percentage of the data


$$Gini = 1 - \sum_i^C p_i^2$$



Decision tree in R – rpart

- ▶ **rpart** is an R-package to fit decision trees (impurity measure based)
- ▶ Here an example of fitting a (classification) tree to the "Iris" data set:

```
library(rpart)
tree.iris <- rpart(Species ~., data=iris) # default uses Gini-index
plot(tree.iris, margin = 0.1)
text(tree.iris, use.n = T)
# Predict training data (use type = 'prob' to get probabilities):
pred.tree <- predict(tree.iris, iris, type = 'class')
# Confusion matrix:
(confT <- table(pred.tree, iris$Species))
pred.tree   setosa versicolor virginica
setosa      50         0         0
versicolor  0         49         5
virginica    0         1        45
# Training-error:
diag(confT) <- 0
missCount <- sum(confT)
(trainErr <- missCount/nrow(iris))
[1] 0.04
```




Decision tree in R – rpart

- **rpart** is an R-package to fit decision trees (impurity measure based)
- Here an example of fitting a (classification) tree to the "Iris" data set:


```
library(rpart)
tree.iris <- rpart(Species ~., data=iris) # default uses Gini-index
plot(tree.iris, margin = 0.1)
text(tree.iris, use.n = T)
# Predict training data (use type = 'prob' to get probabilities):
pred.tree <- predict(tree.iris, iris, type = 'class')
# Confusion matrix:
(confT <- table(pred.tree, iris$Species))
pred.tree   setosa versicolor virginica
  setosa      50         0         0
versicolor   0         49         5
virginica     0          1        45
# Training-error:
diag(confT) <- 0
missCount <- sum(confT)
(trainErr <- missCount/nrow(iris))
[1] 0.04
```

What is this?



Problems with rpart

- ▶ Using an impurity measure based approach to select splits suffers from an **inherent bias** in certain cases
 - ▶ This bias is related to the **number of splits** that are possible for each variable
 - ▶ How many splits are there per variable?
 - ▶ For **continuous** variable (numeric):



- ▶ For the five data points above, there are only **four meaningful** splits possible (along X_1)
- ▶ What about **other types** of variables?

Problems with rpart

- ▶ We can also include **categorical variables** (factors) as predictors in a decision tree

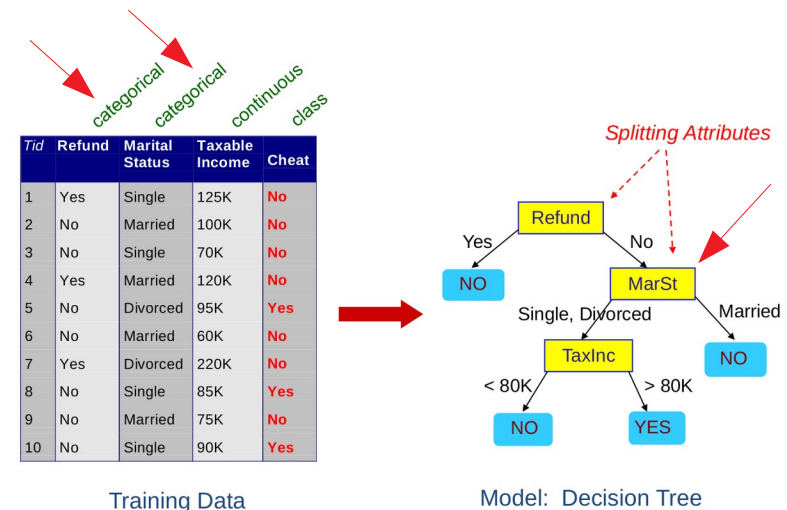
- ▶ How many splits are possible for factors?

- ▶ E.g. **two-level** factor (levels: a,b)
a | b **split1**

- ▶ E.g. **four-level** factor (levels: a,b,c,d):
a | b,c,d **split1**
b | a,c,d **split2**
c | a,b,d **split3**
d | a,b,c **split4**
a,b | c,d **split5**
a,c | b,d **split6**
a,d | b,c **split7**

- ▶ **k-level** factor: $\boxed{\text{number of splits} = 2^{(k-1)} - 1}$ → 10-levels: **511 splits**

- ▶ With increasing number of levels we can apply an **exponentially growing** number of splits



Source: Tan, Steinbach, Kumar

Problems with rpart


- ▶ With increasing number of levels we can apply an **exponentially growing** number of splits
- ▶ The difference in numbers of splits between continuous variables, factors with few levels and factors with many levels creates an issue of **multiple comparisons**
- ▶ Multiple comparisons problem in decision trees:
 - ▶ Even if there is **no real relation** between predictors and the target variable, under certain circumstances multi-level factors will be preferably selected for splitting compared to numerical variables or factors with few levels because they get **more "chances"** to reach a high impurity-decrease (even if this decrease is only achieved by pure luck)
 - ▶ More chances in the sense that more splits can be evaluated (more chances to get lucky)
 - ▶ Ideally, we **do not want** our decision trees to be biased in their variable selection
- ▶ We will see how this bias manifests itself by running a small **simulation study** (Exercise)

E.g. **two**-level factor (levels: a,b)
a | b **split1**

E.g. **four**-level factor (levels: a,b,c,d):
a | b,c,d **split1**
b | a,c,d **split2**
c | a,b,d **split3**
d | a,b,c **split4**
ab | c,d **split5**
ac | b,d **split6**
ad | b,c **split7**

Bias-free partitioning

- ▶ **Bias-free partitioning** has been proposed to solve this problem (e.g. see Hothorn et al. 2006)
 - ▶ Main idea: Use significance tests instead of impurity measures to select variable and next split
- ▶ How does it work?
 - ▶ The selection process is more clearly separated into
 - 1) Choosing the next splitting **variable** and
 - 2) Choosing the splitting **location** along the chosen variable



Bias-free partitioning

- ▶ Choosing the next splitting variable:
 - ▶ Test for **each variable** whether it has a **significant association** with the target variable (e.g. think of correlation test for numerical variables and of ANOVA for categorical variables). Collect the associated p-values (can be compared between tests)
 - ▶ If no variable shows a significant p-value stop splitting (integrated **stopping rule**)
 - ▶ Otherwise choose the variable with the **lowest p-value** for the next split

Bias-free partitioning

► Choosing the next splitting variable:

- Test for **each variable** whether it has a **significant association** with the target variable (e.g. think of correlation test for numerical variables and of ANOVA for categorical variables). Collect the associated p-values (can be compared between tests)
- If no variable shows a significant p-value stop splitting (integrated **stopping rule**)
- Otherwise choose the variable with the **lowest p-value** for the next split

Actual implementation uses special permutation tests (called "**conditional inference tests**") that are available for all types of variables.

Bias-free partitioning

Actual implementation uses special permutation tests (called "**conditional inference tests**") that are available for all types of variables.

► Choosing the next splitting variable:

- Test for **each variable** whether it has a **significant association** with the target variable (e.g. think of correlation test for numerical variables and of ANOVA for categorical variables). Collect the associated p-values (can be compared between tests)
- If no variable shows a significant p-value stop splitting (integrated **stopping rule**)
- Otherwise choose the variable with the **lowest p-value** for the next split

► Choosing the next splitting location:

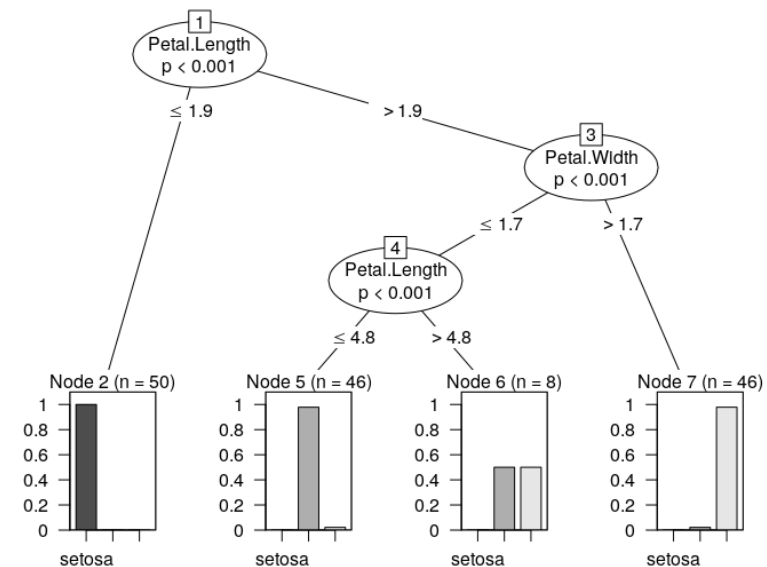
- Calculate for each possible split (along the chosen variable) a **test-statistic** (expressing the difference between the created partitions, *could also use impurity measure*) and select split with the highest test-statistic
-
- Using this approach, there is **no bias** in the selection of different variable-types

Bias-free partitioning in R

- "party" is an R-package which includes functions to fit bias-free (significance test based) decision trees (**ctree()**-function)
- Fitting a classification tree to the "Iris" data set

```
library(party)
ctree.iris <- ctree(Species ~., data=iris) # fit tree
plot(ctree.iris)
pred.ctree <- predict(ctree.iris, newdata=iris, type='response')
(confT <- table(pred.ctree, iris$Species)) # Confusion matrix
```

pred.ctree	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	49	5
virginica	0	1	45




Decision trees for regression

- ▶ Decision trees can be used for classification (categorical target variable) and for **regression** (numerical target variable)
- ▶ The general principle behind building the tree structure does not change
 - ▶ E.g. use **mean squared error** as "impurity measure" instead of Gini-index

$$MSE = \frac{1}{n} \sum_i^n (x_i - \bar{x})^2$$

- ▶ We try to find the splits which best reduce the mean squared error
- ▶ With **ctree** we can fit a bias-free regression tree (again based on significance tests)




Decision trees for regression

- ▶ Decision trees can be used for classification (categorical target variable) and for **regression** (numerical target variable)
- ▶ The general principle behind building the tree structure does not change
 - ▶ E.g. use **mean squared error** as "impurity measure" instead of Gini-index

$$MSE = \frac{1}{n} \sum_i^n (x_i - \bar{x})^2$$

- ▶ We try to find the splits which best reduce the mean squared error
- ▶ With **ctree** we can fit a bias-free regression tree (again based on significance tests)



Regression tree in R

- ▶ Example of fitting a regression tree to the "Iris" data using `ctree()`

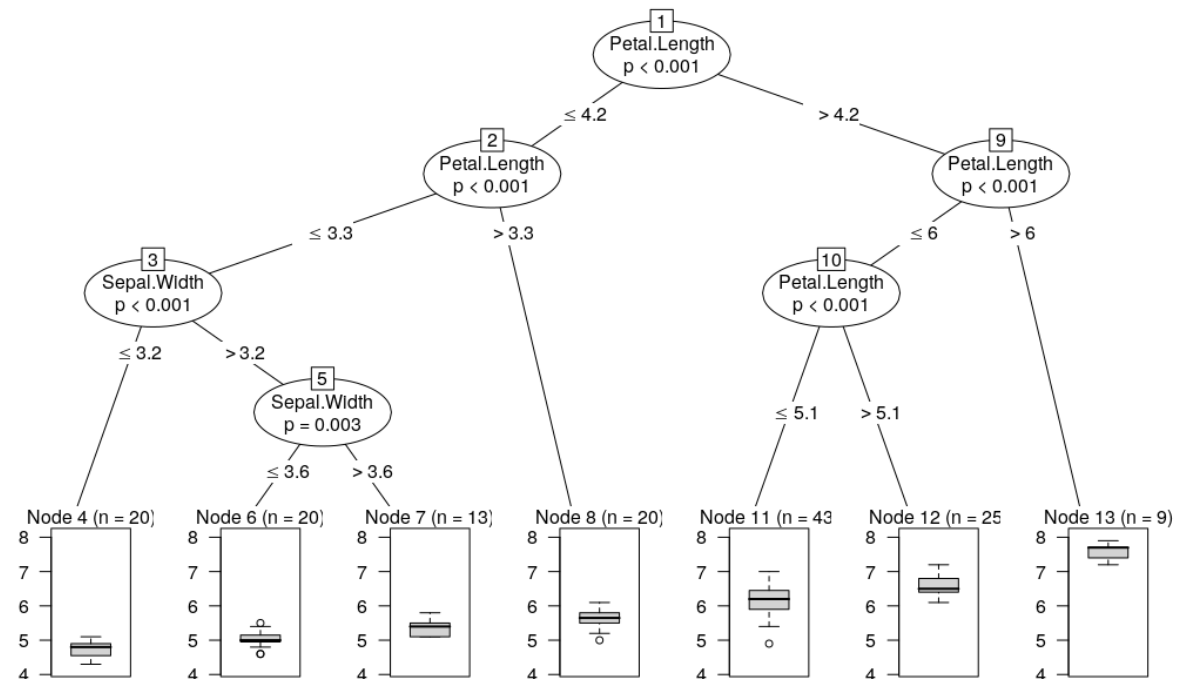
```
library(party)
```

```
ctree.iris <- ctree(Sepal.Length ~., data=iris) # fit regression tree
```

```
plot(ctree.iris)
```

```
# Get the predicted values (for training data):
```

```
pred.ctree <- predict(ctree.iris, iris, type='response')
```



Decision trees – Summary

► Advantages of decision trees:

- Good interpretability
- No assumptions regarding distributions
- Robust to outliers
- Can capture non-linear structures and complex interactions
- Low bias in prediction with sufficient depth

► Disadvantages:

- Single trees are instable (small changes in data can give different-looking tree, especially if underlying pattern is complex)
- In case of not pruned rpart-trees: Tend to overfit
- Needs a lot of data to capture linear structures

Caution: In case of different types of predictor variables use algorithm with bias-free variable selection (ctree in R)

Further reading

- ▶ Strobl, C., Malley, J., & Tutz, G. (2009). **An introduction to recursive partitioning: Rationale, application, and characteristics of classification and regression trees, bagging, and Random Forests.** *Psychological Methods*, 14 (4), 323–348. doi: 10.1037/a0016973