


R-course:
Machine Learning using R

Random Forest




Yannick Rothacher

Zürich, 2021


Recap: Decision trees

- ▶ Last lecture we saw how decision trees can be used for **classification** and **regression**
- ▶ Looked at the different methods for **variable** and **split selection** in `rpart()` and `ctree()` function
- ▶ One advantage of decision trees is their **good interpretability**
- ▶ **However:** If trees are allowed to grow deep they show large variability and low bias (overfitting the data)
- ▶ What can be done to prevent overfitting?



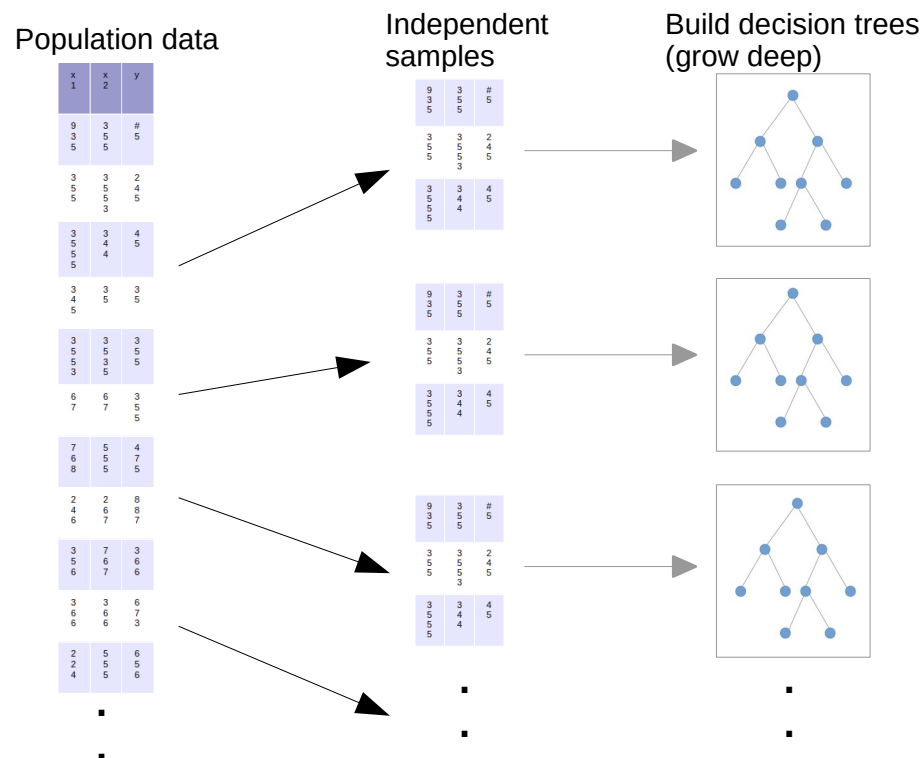
Ensemble methods

- ▶ Why not combine multiple classifiers to improve prediction?
- ▶ Given some training data and a test-observation:
 - ▶ You could apply multiple, different classifiers to the data (logistic regression, **KNN**, **decision tree**, neural network, ...) and incorporate each classifier's prediction into your final prediction of the test observation (e.g. majority vote)
- ▶ **Ensemble methods** are based on a similar intuition: Combine multiple (simple) classifiers of the same type to improve performance
- ▶ With regard to decision trees there are two ensemble methods, which are mostly applied:
 - ▶ Bagging
 - ▶ Random Forest



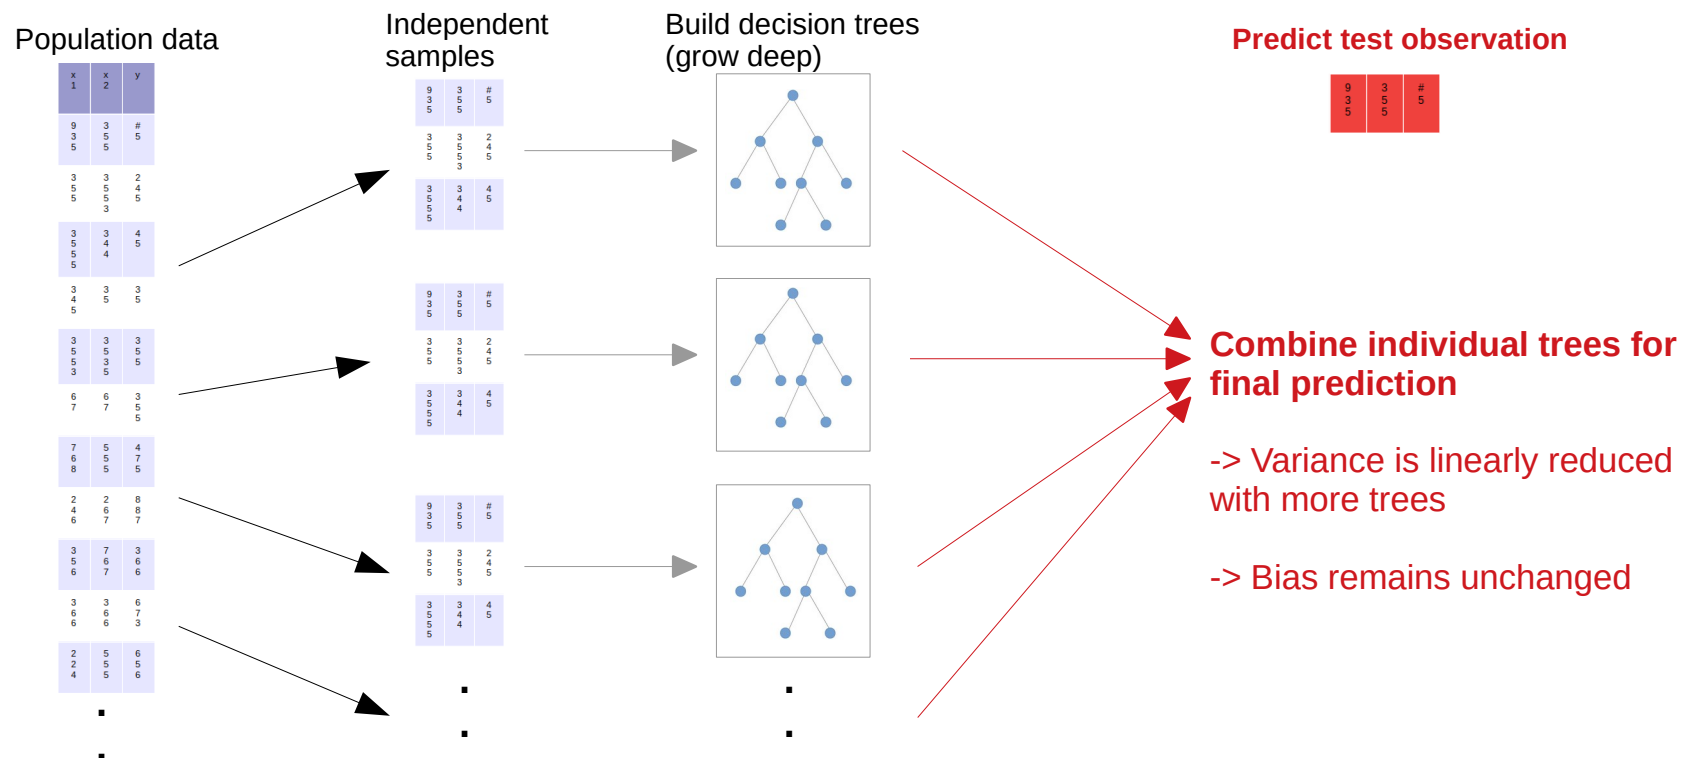
Ensemble methods: Bagging

- ▶ **Bagging** (bootstrap aggregating) is an ensemble method typically applied to decision trees (but can be applied to all classifiers)
- ▶ Main idea: Fully grown decision trees have **low bias** but **high variance**...
- ▶ ...it would be nice to have multiple independent samples from the same population to **reduce variance** by taking the mean of prediction




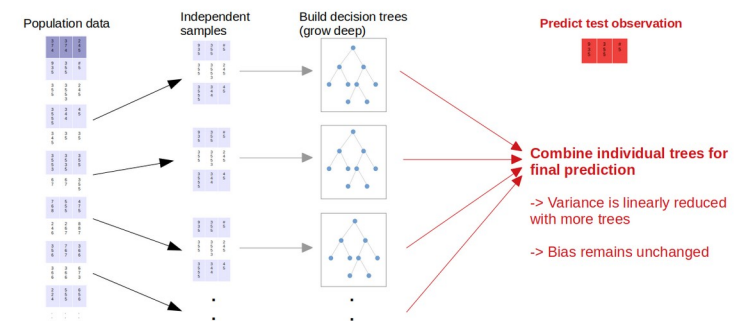
Ensemble methods: Bagging

- ▶ **Bagging** (bootstrap aggregating) is an ensemble method typically applied to decision trees (but can be applied to all classifiers)
- ▶ Main idea: Fully grown decision trees have **low bias** but **high variance**...
- ▶ ...it would be nice to have multiple independent samples from the same population to **reduce variance** by taking the mean of prediction




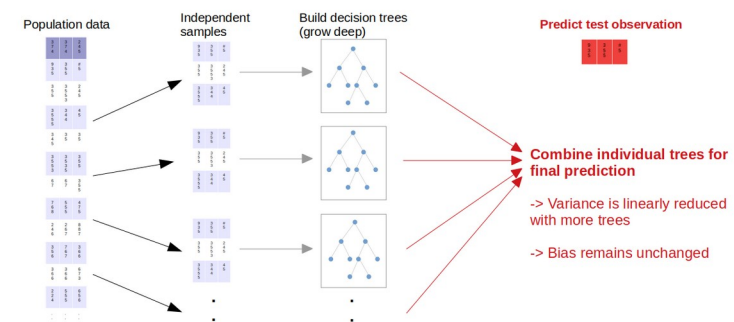
Ensemble methods: Bagging

- ▶ **Bagging** (bootstrap aggregating) is an ensemble method typically applied to decision trees (but can be applied to all classifiers)
- ▶ Main idea: Fully grown decision trees have **low bias** but **high variance**...
 - ▶ ...it would be nice to have multiple independent samples from the same population to **reduce variance** by taking the mean of prediction
- ▶ By averaging the individual trees the variance of the estimation is reduced



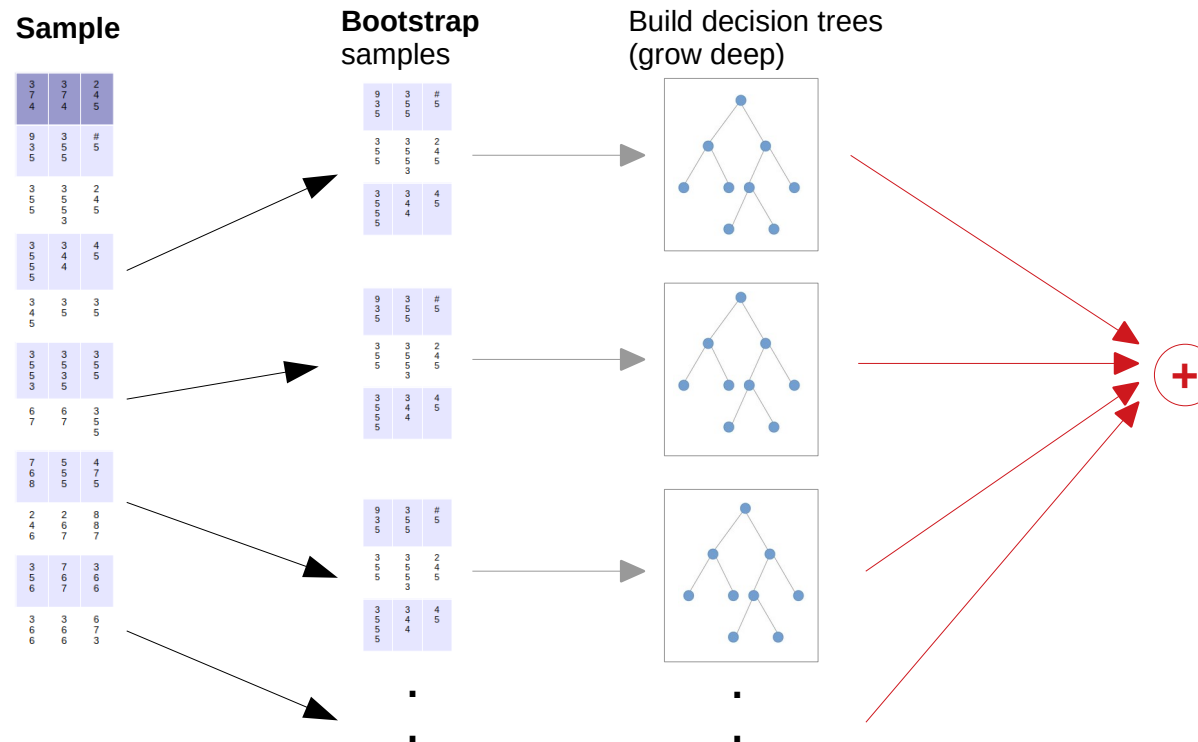
Ensemble methods: Bagging

- ▶ **Bagging** (bootstrap aggregating) is an ensemble method typically applied to decision trees (but can be applied to all classifiers)
- ▶ Main idea: Fully grown decision trees have **low bias** but **high variance**...
 - ▶ ...it would be nice to have multiple independent samples from the same population to **reduce variance** by taking the mean of prediction
- ▶ By averaging the individual trees the variance of the estimation is reduced
- ▶ **However:** In reality we usually only have **one training sample** from the population



Ensemble methods: Bagging

- ▶ In reality we usually only have only **one sample** from the population
 - ▶ Solution: We now take **bootstrap samples** from our training data
 - ▶ Grow deep decision trees on bootstrap samples and combine the trees for prediction



Ensemble methods: Bagging

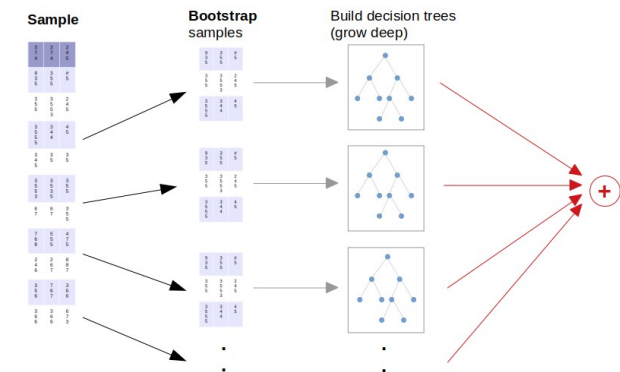
- ▶ In reality we usually only have only **one sample** from the population
 - ▶ Solution: We now take **bootstrap samples** from our training data (originally suggested by Breiman 1996)
 - ▶ Grow deep decision trees on bootstrap samples and combine the trees for prediction
- ▶ Bootstrap samples are samples drawn with replacement:

Original data

4	5	22	3	99	67	43	1	21	2
---	---	----	---	----	----	----	---	----	---

Bootstrap samples (same size as original data)

3	5	67	5	43	67	21	2	99	3
2	4	4	67	21	1	1	67	99	3
5	5	21	99	1	2	22	99	67	5



Ensemble methods: Bagging

- ▶ In reality we usually only have only **one sample** from the population
 - ▶ Solution: We now take **bootstrap samples** from our training data (originally suggested by Breiman 1996)
 - ▶ Grow deep decision trees on bootstrap samples and combine the trees for prediction
- ▶ Bootstrap samples are samples drawn with replacement:

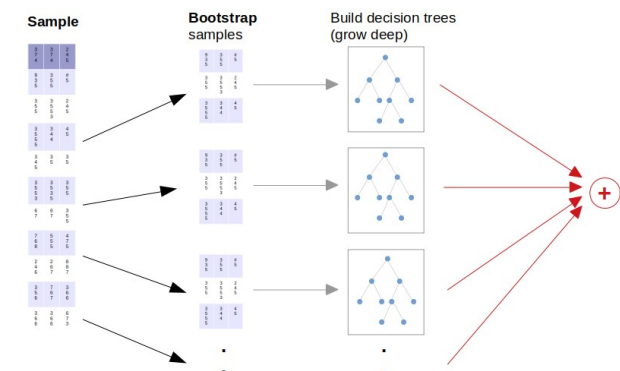
Original data

4	5	22	3	99	67	43	1	21	2
---	---	----	---	----	----	----	---	----	---

Bootstrap samples (same size as original data)

3	5	67	5	43	67	21	2	99	3
2	4	4	67	21	1	1	67	99	3
5	5	21	99	1	2	22	99	67	5

Same observation can appear multiple times!



Ensemble methods: Bagging

- ▶ In reality we usually only have only **one sample** from the population
 - ▶ Solution: We now take **bootstrap samples** from our training data (originally suggested by Breiman 1996)
 - ▶ Grow deep decision trees on bootstrap samples and combine the trees for prediction
- ▶ Bootstrap samples are samples drawn with replacement:


Original data

4	5	22	3	99	67	43	1	21	2
---	---	----	---	----	----	----	---	----	---

Bootstrap samples (same size as original data)

3	5	67	5	43	67	21	2	99	3
2	4	4	67	21	1	1	67	99	3
5	5	21	99	1	2	22	99	67	5

Same observation can appear multiple times!




Bagging shows better performance than one decision tree applied to original data (why?)

Note: Individual bootstrap samples are not independent!


-> Variance is only reduced sublinearly

Ensemble methods: Bagging

Decision Boundary: Tree



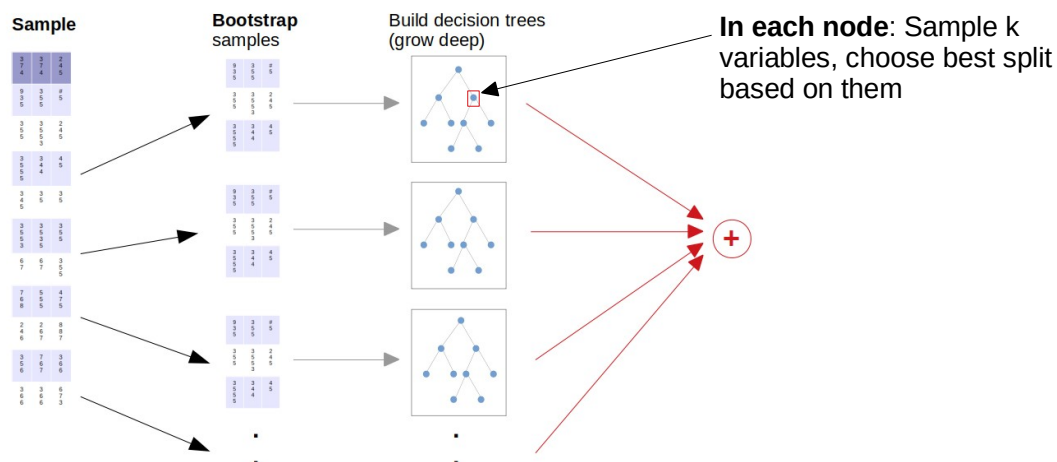
Decision Boundary: Bagging



Quelle: Ji Zhu, University of Michigan

Ensemble methods: Random Forest

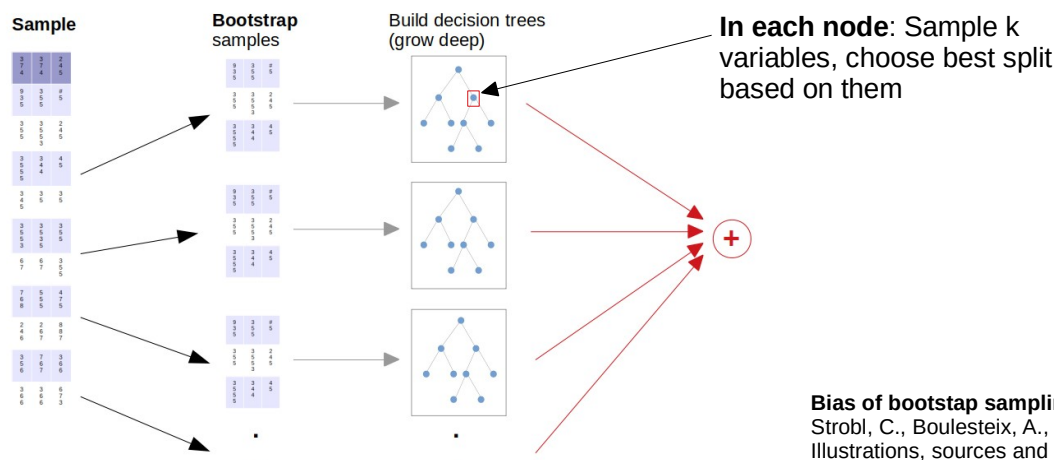
- ▶ **Random forest** is an extension of bagging
 - ▶ Main idea: Do not only sample from the data but also from the **variables**, which are used for splitting
- ▶ In random forest the trees are generated in the following way
 - ▶ Generate bootstrap samples from original data (same like bagging)
 - ▶ Build a decision tree on each bootstrap sample, but...
 - ▶ ... at each node of a decision tree, **randomly select k variables**, which are evaluated for splitting. Choose the best split (using only the k variables).



Ensemble methods: Random Forest

- ▶ **Random forest** is an extension of bagging
 - ▶ Main idea: Do not only sample from the data but also from the **variables**, which are used for splitting
- ▶ In random forest the trees are generated in the following way
 - ▶ Generate bootstrap samples from original data (same like bagging)
 - ▶ Build a decision tree on each bootstrap sample, but...
 - ▶ ... at each node of a decision tree, **randomly select k variables**, which are evaluated for splitting. Choose the best split (using only the k variables).

Can also use **sub-samples** instead of bootstrap samples. Bootstrap sampling can also induce bias in variable selection.

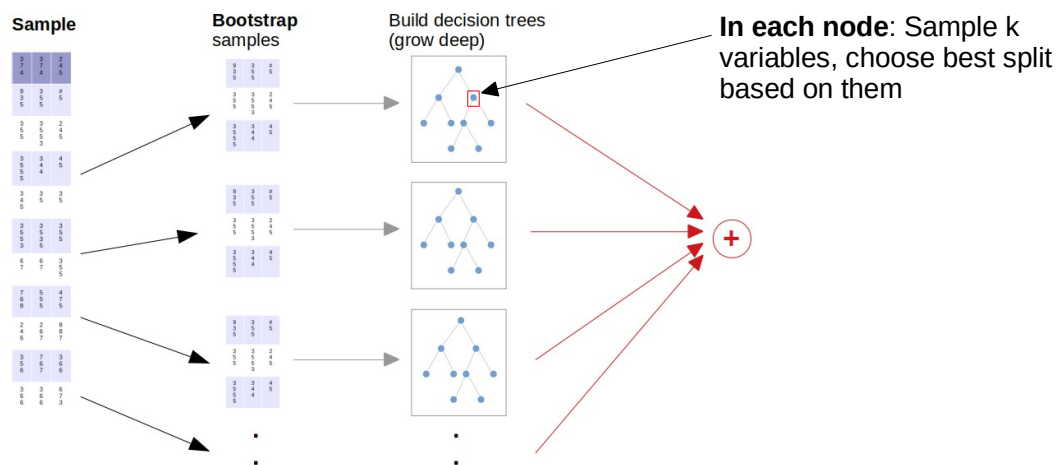


Bias of bootstrap sampling in RF:

Strobl, C., Boulesteix, A., Zeileis, A. et al. Bias in random forest variable importance measures: Illustrations, sources and a solution. BMC Bioinformatics 8, 25 (2007)


Ensemble methods: Random Forest

- ▶ What is the point of sampling the variables?
 - ▶ Sampling variables de-correlates the individual trees, makes them more diverse (this increases variance reduction)
 - ▶ **Intuition:** By restricting the variable selection at each node, some variables are incorporated in the analysis which might otherwise never be considered
 - ▶ This can reveal interactions in the data, which would otherwise not be detected



Ensemble methods: Random Forest

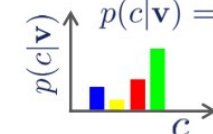
- ▶ Prediction of a new observation same like in bagging:
 - ▶ Run observation through all trees and incorporate each tree's prediction in the final prediction (e.g. majority vote for classification or mean for regression)
- ▶ Example of aggregating the tree's results for **classification**:



Source: Microsoft

Winner category: Category that was predicted most often (green)

Average probability:

$$p(c|\mathbf{v}) = \frac{1}{T} \sum_t^T p_t(c|\mathbf{v})$$


Random Forest in R

- We will use the **cforest()** function from the package "party" (unbiased variable selection)

```
library(party)
cfor_ctr <- cforest_unbiased(ntree = 500, mtry = 2)
(rf.iris <- cforest(Species~., data = iris, controls = cfor_ctr))
```



Random Forest using Conditional Inference Trees

Number of trees: 500

Response: Species

Inputs: Sepal.Length, Sepal.Width, Petal.Length, Petal.Width

Number of observations: 150

honest cross classification (OOB-error)

```
(confT <- table(iris$Species, predict(rf.iris, OOB = TRUE)))
```

	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	47	3
virginica	0	4	46

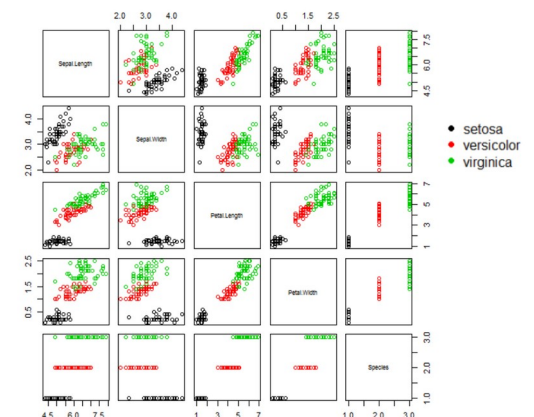
```
diag(confT) <- 0
```

```
sum(confT)/nrow(iris)
```

```
[1] 0.0466667
```

```
varimp(rf.iris)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
0.051127273	0.002109091	0.314800000	0.264909091



Random Forest in R

- We will use the **cforest()** function from the package "party" (unbiased variable selection)

```
library(party)
cfor_ctr <- cforest_unbiased(ntree = 500, mtry = 2)
(rf.iris <- cforest(Species~., data = iris, controls = cfor_ctr))
```

Define the options/parameters for random forest

ntree: How many trees are grown?

mtry: How many variables are sampled at each node?

Random Forest using Conditional Inference Trees

Number of trees: 500

Response: Species

Inputs: Sepal.Length, Sepal.Width, Petal.Length, Petal.Width

Number of observations: 150

honest cross classification (OOB-error)

```
(confT <- table(iris$Species, predict(rf.iris, OOB = TRUE)))
```

What is OOB? ...

setosa versicolor virginica

setosa	50	0	0
versicolor	0	47	3
virginica	0	4	46

```
diag(confT) <- 0
```

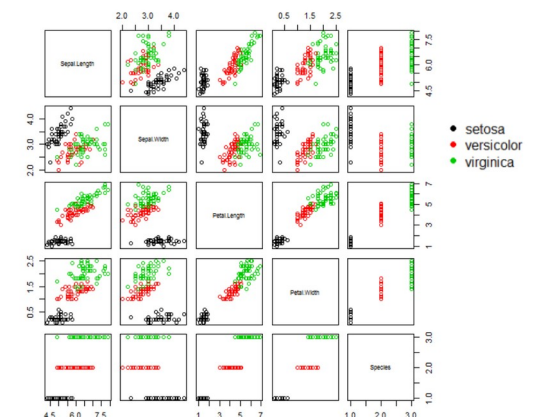
```
sum(confT)/nrow(iris)
```

```
[1] 0.0466667
```

```
varimp(rf.iris)
```

What is varimp? ...

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
0.05112723	0.002109091	0.314800000	0.264909091



Out-of-bag error (OOB)




- ▶ Random Forest comes with its own integrated evaluation tool!
- ▶ Each tree is fitted to a bootstrap sample (or sub-sample) of original data
 - ▶ Thus, every tree in the forest has only seen a part of the data
- ▶ To calculate the **OOB-error**:
 - ▶ Predict each observation in the data using only the trees, which have not seen the observation when they were generated

Out-of-bag error (OOB)


- ▶ Random Forest comes with its own integrated evaluation tool!
- ▶ Each tree is fitted to a bootstrap sample (or sub-sample) of original data
 - ▶ Thus, every tree in the forest has only seen a part of the data
- ▶ To calculate the **OOB-error**:
 - ▶ Predict each observation in the data using only the trees, which have not seen the observation when they were generated

e.g. observation #24:

x1	x2	x3	y
23.4	66.3	2	A

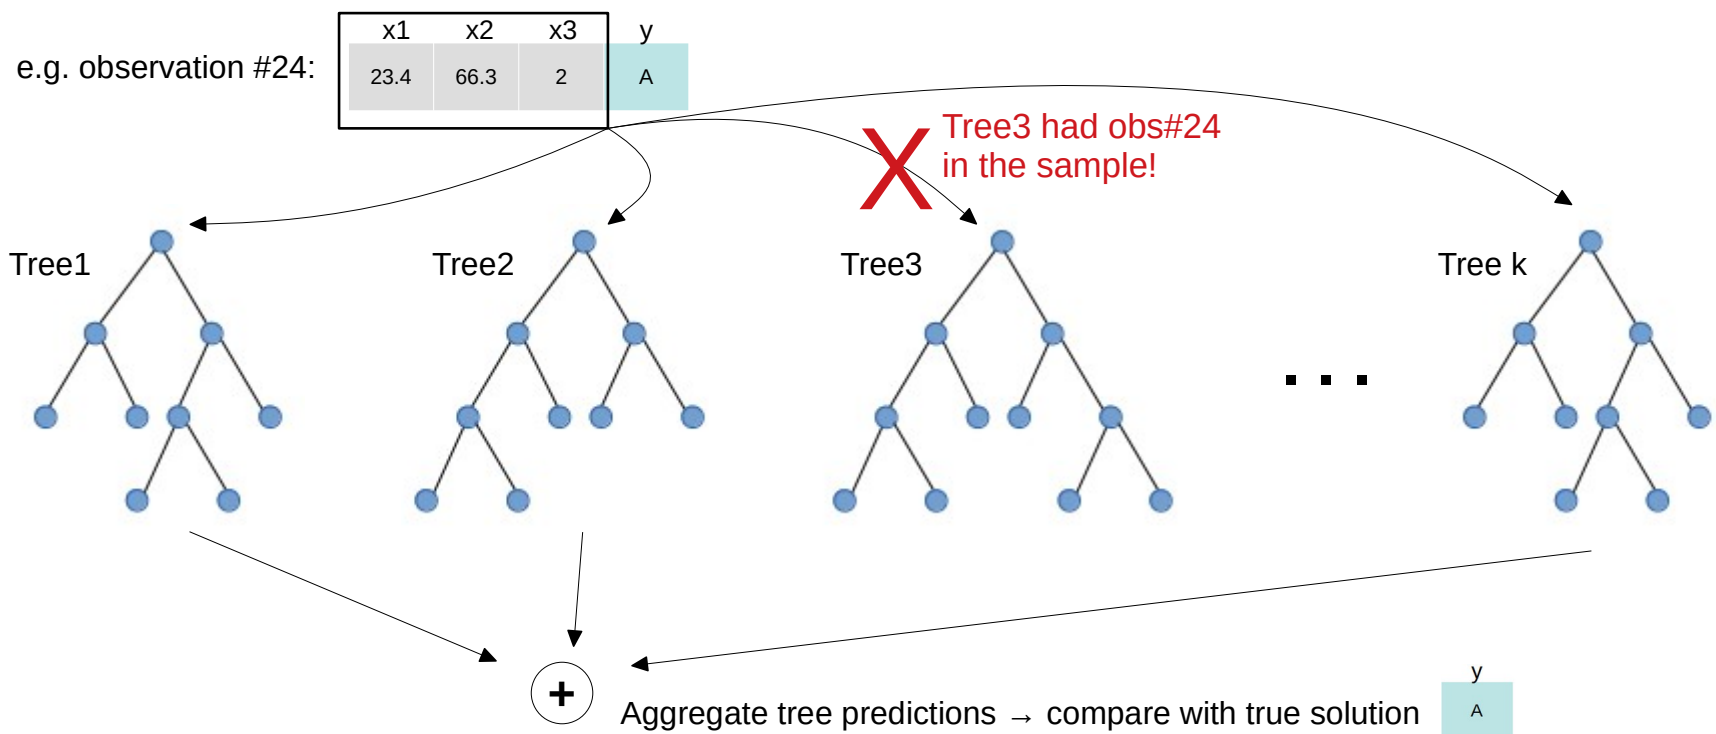


...



Out-of-bag error (OOB)

- ▶ Random Forest comes with its own integrated evaluation tool!
- ▶ Each tree is fitted to a bootstrap sample (or sub-sample) of original data
 - ▶ Thus, every tree in the forest has only seen a part of the data
- ▶ To calculate the **OOB-error**:
 - ▶ Predict each observation in the data using only the trees, which have not seen the observation when they were generated




Out-of-bag error (OOB)

- ▶ Create OOB confusion table to look at the OOB-performance (see R-code slide)
- ▶ Use `predict()`-function and set `OOB=TRUE`, give no "newdata" argument

```
rf.iris <- cforest(Species~., data = iris, controls =  
cfor_ctr)  
(confT <- table(iris$Species, predict(rf.iris, OOB = TRUE)))  
      setosa versicolor virginica  
setosa      50         0         0  
versicolor   0        47         3  
virginica     0         4        46  
diag(confT) <- 0  
sum(confT)/nrow(iris)  
[1] 0.04666667
```

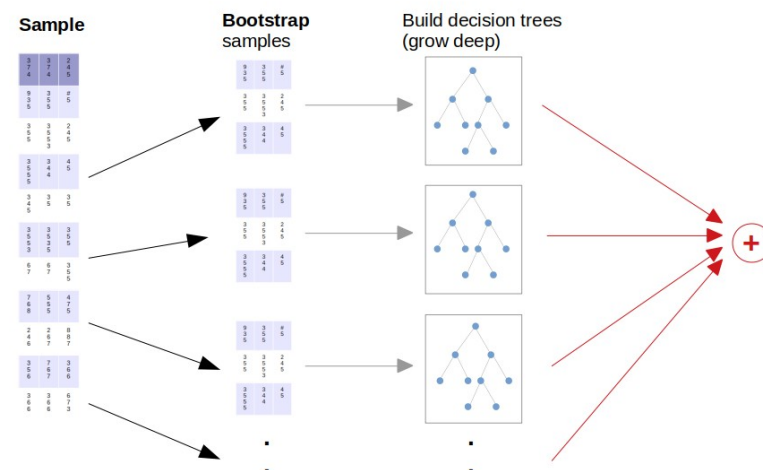
Interpretability of Random Forest

- ▶ One advantage of decision trees are their good interpretability
 - ▶ Can easily track how each variable affects the prediction
- ▶ In Random Forest we somewhat lose the interpretability
 - ▶ Difficult to track how each variable is active in each of the e.g. 500 trees...
- ▶ How can we assess which variables are **important** for prediction?




Model: Decision Tree

VS.



Variable importance in RF


- ▶ How could variable importance be **scored**?
 - ▶ Simple idea: Count the number of times a variable is chosen for splitting throughout all trees
- ▶ Two more elaborate approaches:
- ▶ **Variable importance 1**: Average score improvement at splits
 - ▶ For each node where variable x was used for splitting, record the achieved improvement in the splitting score (e.g. Gini-index -> **Gini-importance score**)
 - ▶ Variable x 's importance 1 score is the average of all score improvements (can be weighted with the number of data points in each node)



Variable importance in RF

- ▶ How could variable importance be **scored**?
 - ▶ Simple idea: Count the number of times a variable is chosen for splitting throughout all trees
- ▶ Two more elaborate approaches:
 - ▶ **Variable importance 1**: Average score improvement at splits
 - ▶ For each node where variable x was used for splitting, record the achieved improvement in the splitting score (e.g. Gini-index -> **Gini-importance score**)
 - ▶ Variable x 's importance 1 score is the average of all score improvements (can be weighted with the number of data points in each node)

Caution: Gini-Importance is biased if different types of variables are present



Variable importance in RF

- ▶ **Variable importance 2: Permutation importance**
- ▶ **Permutation importance** is a very intuitive importance score
- ▶ Main idea: Mix up the values of variable x to break up any meaningful relation between x and the target variable (permutation)
- ▶ Check how much the performance (usually the OOB-error) drops after permutation of x

Original data	Predictor 1	Predictor 2	Predictor 3	Target var
	X1	X2	X3	Y
	2.3	33.1	67	A
	5.1	35.8	70.3	B
	3.3	34.0	96	A
	2.8	37.7	85	C
	1.3	38.3	84.9	A

VS.

X1 permuted	Predictor 1	Predictor 2	Predictor 3	Target var
	X1	X2	X3	Y
	5.1	33.1	67	A
	2.3	35.8	70.3	B
	2.8	34.0	96	A
	1.3	37.7	85	C
	3.3	38.3	84.9	A

→ Calculate **OOB-error**

→ Calculate **OOB-error**

Permutation importance:
Difference in OOB-error

With stronger relation between x1 and target variable, the more should the OOB-error grow

Variable importance with R

- To calculate the variable importance after fitting a Random Forest use the **varimp()**-function (default is permutation importance score)


```
rf.iris <- cforest(Species~., data = iris, controls = cfor_ctr)
varimp(rf.iris)
```

```
Sepal.Length Sepal.Width Petal.Length Petal.Width
0.052727273 0.001854545 0.305636364 0.272727273
```

```
# Function to plot importance scores nicely
```

```
implot <- function(rf.fit){
  par(mar=c(9,9,9,9))
  scores <- sort(varimp(rf.fit))
  plot(x = scores, y = 1:length(scores) ,
       xlim=c(-max(scores), max(scores)*1.1),
       ylim=c(-1,(length(scores)+1)),
       yaxt='n', ylab='', xlab='Importance score')
  axis(2, at=1:length(scores), labels = names(scores), las=1)
  abline(h=1:length(scores), lty=2, col='grey')
  abline(v=0, col='red')
}
```

```
implot(rf.iris)
```



Variable importance with R

- ▶ To calculate the variable importance after fitting a Random Forest use the **varimp()**-function (default is permutation importance score)

```
rf.iris <- cforest(Species~., data = iris, controls = cfor_ctr)
varimp(rf.iris)
```


```
Sepal.Length Sepal.Width Petal.Length Petal.Width
0.052727273 0.001854545 0.305636364 0.272727273
```

```
# Function to plot importance scores nicely
```

```
implot <- function(rf.fit){
  par(mar=c(9,9,9,9))
  scores <- sort(varimp(rf.fit))
  plot(x = scores, y = 1:length(scores) ,
       xlim=c(-max(scores), max(scores)*1.1),
       ylim=c(-1,(length(scores)+1)),
       yaxt='n', ylab='', xlab='Importance score')
  axis(2, at=1:length(scores), labels = names(scores), las=1)
  abline(h=1:length(scores), lty=2, col='grey')
  abline(v=0, col='red')
}
```

```
implot(rf.iris)
```

Permutation importance
scores can also be negative!
→ Why?




Because importance scores are based on random permutation we should fix the random seed when calculating them

Dependency plots in Random Forest

- ▶ Dependency plots are another method to "bring light" into what is happening in the forest
- ▶ Main idea: See how the prediction of the target variable changes when only **one predictor** is shifted in its value
- ▶ Example of **partial dependency plot** on RF-regression on the iris data (target variable: Sepal.width):


```
cfor_ctr <- cforest_unbiased(ntree = 500, mtry = 2)
rf.SepWid <- cforest(Sepal.Width~., data = iris, controls = cfor_ctr)
steps <- seq(min(iris$Sepal.Length), max(iris$Sepal.Length), by = 0.1)
predic <- matrix(NA, nrow=nrow(iris), ncol = length(steps))
for (i in 1:nrow(iris)){
  obs <- iris[i,]
  obs$Sepal.Width <- NULL # Remove column
  for(s in 1:length(steps)){
    obs$Sepal.Length <- steps[s]
    predic[i,s] <- predict(rf.SepWid, newdata=obs)
  }
}
plot(NULL, xlim=c(min(iris$Sepal.Length), max(iris$Sepal.Length)),
     ylim=c(min(iris$Sepal.Width), max(iris$Sepal.Width)),
     xlab='Sepal.Length', ylab='Predicted Sepal.Width')
for(l in 1:nrow(predic)){
  lines(x = steps, y = predic[l,], col='grey')
}
lines(x=steps, y=apply(predic, 2, mean), col='red', lwd=2)
```



Partial dependency plots with "pdp"

- ▶ **"pdp"** is an R package to construct partial dependency plots
- ▶ Can be applied to different Machine Learning methods
 - ▶ Also works for Random Forests fitted with `cforest()`

```
cfor_ctr <- cforest_unbiased(ntree = 500, mtry = 2)
rf.SepWid <- cforest(Sepal.Width~., data = iris,
                     controls = cfor_ctr)
#install.packages("pdp")
library(pdp)
partial(rf.SepWid, pred.var = "Sepal.Length", plot = TRUE)
```



- ▶ Meaning of curve (same like in previous slide):
 - ▶ Mean predicted Sepal.Width for varying Sepal.Length (averaged over all observed constellations of other co-predictors)

Summary Random Forest

► Advantages:

- Random Forest is a very effective ML method (good performance in ML-competitions)
- Is non-parametric, poses no assumption regarding distribution of variables or residuals
- Comes with included performance evaluation (OOB-error)
- Random Forest does not tend to overfit
- Can handle thousands of input variables
- Can handle lots of noise variables with only few relevant variables
- We can derive variable importance measures
- Random Forest can detect strong and local interactions
- Is robust against outliers
- It can handle imbalanced data (e.g. using stratified bootstrap sampling)
- ...

Summary Random Forest

► Disadvantages:

- Less interpretability, feeling of a "black box"
- Can not extrapolate predictions beyond the range of training data well
- Takes more time to train than e.g. decision tree

Further reading

- ▶ Strobl, C., Malley, J., & Tutz, G. (2009). **An introduction to recursive partitioning: Rationale, application, and characteristics of classification and regression trees, bagging, and Random Forests.** Psychological Methods, 14(4), 323–348.
<https://doi.org/10.1037/a0016973>